

## Exercise 4 - Language Models and Attention

Deadline: 23.01.2023

Total Marks: 30

### Submission

- Submissions through OLAT. Only one group member needs to submit it.
- Your submission should contain a PDF with the solutions to the exercise questions (and a python notebook file) zipped together in a single file.
- Include the group number along with the names and matriculation numbers of all group members on the PDF.
- For the Jupyter notebook, please save them with the outputs of your code displayed.

### 4.1. Language Models

[2 + 2 + 2 + 2 = 8]

Consider a vocabulary  $\mathcal{V} = \{A, B, C\}$  and sequences of length  $T = 10$ .

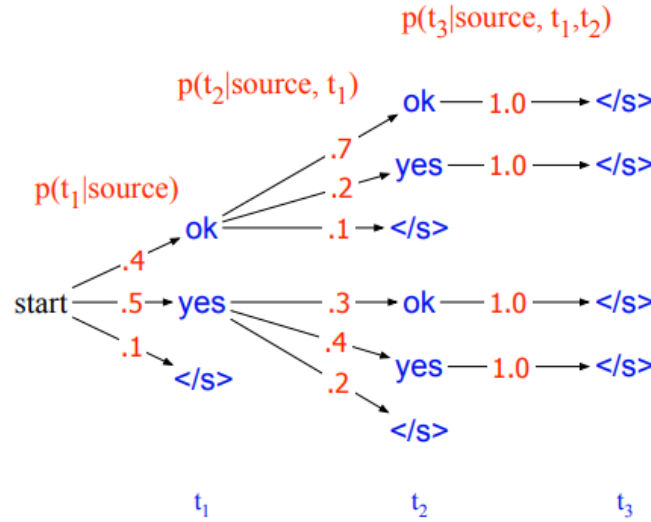
Assume  $p(\mathbf{x}) = \prod_{t=1}^T p(x_t)$  with  $p(x_t = A) = 0.2$ ,  $p(x_t = B) = 0.5$ , and  $p(x_t = C) = 0.3$  for both the model and data distributions.

- Calculate the amount of information (in *bits*) needed to predict the next character in a sequence with a simple (unigram) model.
- Calculate the perplexity  $\text{Perplexity}(p_{data}, p_{model})$  for the same model. Remark on how certain the model is about which letter to predict next.
- Let us now assume that the model distribution is the same as before, but the data distribution is  $p(x_t = A) = 1$  now. Calculate the perplexity again. What does this new perplexity tell you about the model?
- How many bits we need now to encode any possible outcome of  $p_{data}$  using code optimized for  $p_{model}$ ?

## 4.2. Beam Search

[2+5+2=9]

Let  $y = \{yes, ok, </s>\}$  be a vocabulary where  $</s>$  is the end-of-string character. The following figure shows a search tree for generating the target string  $T = t_1, t_2, \dots$  from this vocabulary. Each node represents the conditional probability  $p(t_n | t_1, \dots, t_{n-1}, \mathbf{source})$ , where  $\mathbf{source}$  is the context variable.



- What will be the generated target string in this case with the greedy search algorithm?
- Execute the beam search algorithm with the beam size  $k = 2$ . What is the target string generated in this case? Compute all intermediate probabilities.
- Did the two search algorithms yield the same result? What is the benefit of using beam search over the greedy approach?

## 4.3. Transformers and Attention

[3+5+5=13]

Given

$$X = \begin{bmatrix} 4 & 0 & 2 & 1 \\ 0 & 1 & 1 & 2 \\ 1 & 0 & 1 & 0 \end{bmatrix}, W_Q = \begin{bmatrix} 0 & 0 & 1 \\ 0 & 1 & 0 \\ 0 & 1 & 1 \\ 1 & 0 & 0 \end{bmatrix}, W_K = \begin{bmatrix} 1 & 0 & 1 \\ 1 & 1 & 0 \\ 1 & 1 & 1 \\ 0 & 0 & 1 \end{bmatrix}, W_V = \begin{bmatrix} 0 & 0 & 2 \\ 3 & 0 & 0 \\ 1 & 5 & 0 \\ 0 & 1 & 0 \end{bmatrix}$$

where  $X$  is the input, and  $W_Q, W_K$ , and  $W_V$  are the weights for query, key and value respectively.

- i) Explain the intuition behind the query, key and value in the attention mechanism.
- ii) Compute self-attention for  $X$  showing all the intermediate steps.
- iii) Repeat the same steps you performed in part (ii) using simple PyTorch operations and verify that you get the same answers as in part (ii). Use the provided notebook `Task_4.3.ipynb` for this task.

**Good luck!**